

Henry Mattinson

**Excello: End-user music
programming in Excel**

Computer Science Tripos – Part II

Christ's College

April 6, 2019

Proforma

Name: **Henry Mattinson**
College: **Christ's College**
Project Title: **Excello: End-user music programming in Excel**
Examination: **Computer Science Tripos – Part II, June 2019**
Word Count: **????¹**
Project Originator: **Alan Blackwell**
Supervisor: **Dr Advait Sarkar**

Original Aims of the Project

The main aim of the project was to create a system for music expression and playback allowing users to play individual notes and chords and define their durations, define multiple parts, play loops, define sequences of notes and chords and be able to call these for playback and define the tempo of playback. Followed by the implementation of a converter from an existing musical notation to the Excel system (with compression as an extension) and usability testing of the Excel system.

Work Completed

I designed a notation for music expression in Excel and built a prototype (Excello) satisfying the success criteria above. Participatory design sessions with 21 users served as formative evaluation leading to the implementation of many additional features as extensions. I contributed part of my implementation to an open-source library, this has been merged and published. I built a converter from MIDI to the Excello notation which can convert exactly or perform lossy compression. This was used to translate a corpus of music to the Excello notation. I performed summative evaluation with the users from the participatory design.

¹This word count was computed by summing `texcount chapters/section/content.tex` for each of the five chapters.

Special Difficulties

None.

Declaration

I, Henry Mattinson of Christ's College, being a candidate for Part II of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

Signed [signature]

Date [date]

Contents

1	Introduction	9
1.1	Motivation	9
1.2	Outline of work	10
2	Preparation	11
2.1	Proposal Refinements	11
2.2	Initial Tests	12
2.2.1	Note Synthesis Library	12
2.2.2	Office Javascript API	13
2.3	Excello Design and Language	13
2.3.1	Abstracting Time	13
2.3.2	Initial Prototype Design	14
2.4	Software Engineering	16
2.4.1	Requirements	16
2.4.2	MIDI files	17
3	Implementation	19
3.1	Initial Prototype	19
3.1.1	Turtles	20
3.1.2	Highlighting	21
3.1.3	Chord input	21
3.2	Formative Evaluation	21
3.2.1	Issues and Suggestions	22
3.3	Second Prototype	24
3.3.1	Dynamics	24
3.3.2	Nested Instructions	25
3.3.3	Absolute Tempo	25
3.3.4	Custom Excel Functions	25
3.3.5	Sustain	26
3.3.6	Active Turtles	26
3.3.7	Automatic Movement	26
3.3.8	Inferred Octave	27
3.3.9	Chords	27
3.3.10	Activation of turtles	27

3.4	How it works	28
3.5	MIDI Converter	28
3.6	Repository Overview	28
Bibliography		28

Chapter 1

Introduction

1.1 Motivation

There exist many programs for music notation and composition. Sibelius allows users to write scores using traditional western music notation, whilst music is produced in the live programming interface Sonic Pi by real-time editing of Ruby code [1]. These require users to gain familiarity with a new interface, often with a large threshold to creating simple musical ideas. Spreadsheet users significantly outnumber programmers [15] being the preferred programming language for many people [12]. I believe that this ubiquitousness, along with the affordances of the spreadsheet, would enable new ways to interact with musical notation that capitalise on existing familiarities with spreadsheets and their data handling capabilities.

The use of grid structures is an established concept in music programs, with most sequencing software using one axis of the screen for time and the other for pitch or musical parts. Chris Nash’s Manhattan [11] uses a grid structure where formulae can be defined in the cells to change the cell value, much like in a spreadsheet. However it is limited to columns defining tracks and rows corresponding to different times. Advait Sarkar’s SheetMusic [14] investigated how formulae with sound output can be included within the spreadsheet paradigm. This also introduced abstracting time away from the grid, in this case using an incrementing global `tick` variable which could be referred to in the formulae. Both axes can be used interchangeably for SheetMusic notation or markup that the user wishes to include which is not interpreted musically, a concept idiomatic to Excel usage. Simple formulae such as `if(tick%2==0) p('snare') else p('kick')` allow musical structures to be defined without advanced programming knowledge but quickly become unwieldy for defining larger pieces, especially if they are not highly repetitive. Whilst other spreadsheet music projects exist [4], these simply use the spreadsheet as the medium for conventional sequencing with an auxiliary script used to parse the grid and create musical output.

Excello is an Excel add-in for end-user music programming where users define music in the spreadsheet and can play it back from within Excel. It maintains the abstraction of

time from the grid to keep the flexibility spreadsheets offer but was designed so that the complexity of an individual cell was limited. Existing functionality within Excel can be used, both accelerating the learning curve and increasing the available functionality.

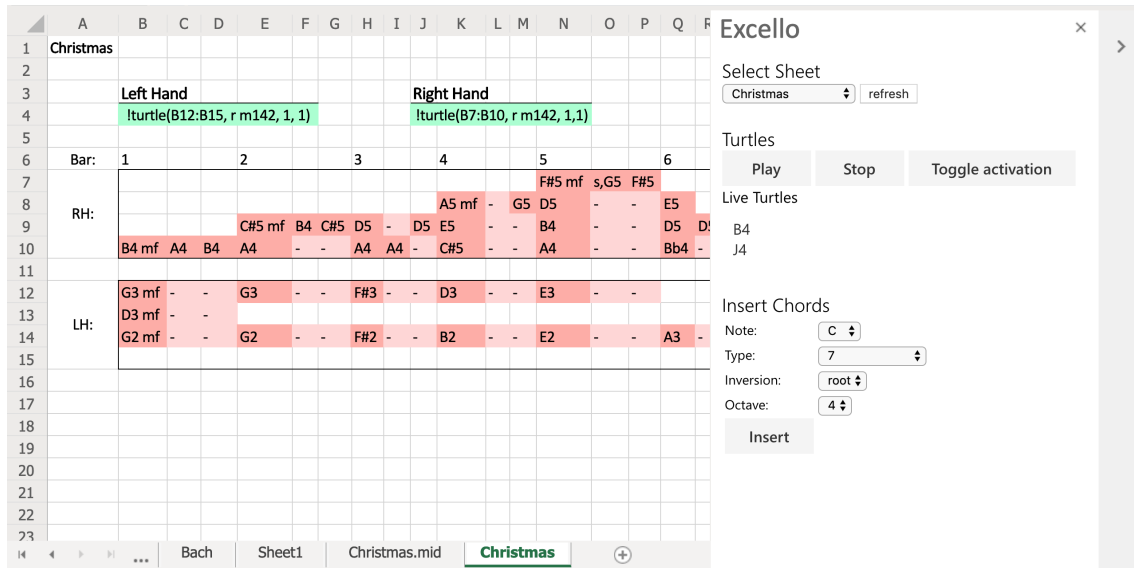


Figure 1.1: The Excello notation and add-in

1.2 Outline of work

1. I design a system for musical expression and playback within Excel. An initial prototype is built satisfying all the success criteria of the system itself: play individual notes and chords with defined durations, define multiple parts, play loops, define sequences of notes and chords and be able to call these for playback, and define the tempo of playback.
2. Participatory design commences using the initial prototype as a discussion point. Following formative evaluation sessions with 21 participants, issues and feature requests are identified. Users continue to use the system and give feedback following their introductory session.
3. A series of additional extensions are implemented to solve the problems identified by participants and add features that are requested to improve the system.
4. A converter from MIDI to the Excello format is built to aid in the translation of a large existing corpus of MIDI files to the Excello notation.
5. Summative evaluation is carried out with the participants. The success of the features implemented as part of the participatory design phase is evaluated. The usability of Excello is analysed using the Cognitive Dimensions of Notation framework [3], focusing on Sibelius as a comparison.

Chapter 2

Preparation

This chapter shall first address the refinements made to the project proposal. Then I shall explain the tests that were performed to establish Excel's suitability for musical development. Next, the design decisions for Excello itself shall be explained. The software engineering tools and techniques employed will then be introduced. Finally, the research that was conducted to decide to implement a converter from MIDI to Excello shall be summarised.

2.1 Proposal Refinements

The project shall invent a notation by which music can be defined within a spreadsheet along with a system for interpreting the notation in the spreadsheet to produce audio output. This shall continue to explore ways in which time can be abstracted away from the grid.

The aim shall be to implement the project as an Excel add-in subject to successful initial testing. An add-in is a web application displayed within Excel that is able to read and write to the spreadsheet using the Office Javascript API. This shall be implemented in a way that allows arbitrary additional data and markup to be included in the spreadsheet. Also, no information beyond the spreadsheet shall need providing for playback via the add-in to be possible. Tests shall be carried out to verify that suitable audio output can be produced for music end-user music programming within Excel to be possible.

A sizeable addition to the project not included in the initial proposal was to perform participatory design [10] to advise on improvements that can be made beyond the initial prototype. The prototype would be introduced to users and from this, new features and improvements implemented. A subset of these participants who gain sufficient familiarity with the project can then be used for more informed summative evaluation. As a result, the proposed extension of incorporating live-coding will only be implemented if there are no other issues or feature requests raised by the userbase that are deemed higher priority.

MIDI shall be the formal notation for which a converter shall be implemented to translate into a CSV file that can then be opened in Excel. Additional explanation on the choice of MIDI is provided below. The choice of MIDI was motivated by participants who wished to be able to integrate Excello in to their use of digital audio workstations such as Logic Pro, Ableton Live and GarageBand.

2.2 Initial Tests

The following section outlines the libraries I explored and the tests carried out to assess the feasibility of synthesising notes given data in a spreadsheet using an Excel add-in. All tests were carried out in Excel Online using Script Lab, an add-in that allows users to create and test simple add-ins experimenting with the Office Javascript API. These add-ins have an HTML front end and can access libraries and data elsewhere online.

A simple add-in that played a wav file stored online was used to verify that an add-in was capable of creating sound.

2.2.1 Note Synthesis Library

The Web Audio API allows audio to be synthesised from the browser using Javascript [9]. To create a program for users to define and play musical structures will require synthesising arbitrary length, pitch and volume notes. In order to avoid the lower-level audio components (e.g. oscillators), I researched libraries that would allow me to deal with higher level musical abstractions of the synthesised notes. Sarkar’s SheetMusic used the library tones¹ which provides a very simple API where only the pitch and volume envelope² of all notes. Other limitations included no definition of volume and only including simple waveform synthesisers.

Tone.js³ is a library built on top of the Web Audio API providing greater functionality than tones. An **Instrument** such as a **Synth** or **Sampler** is defined. The **triggerattackrelease** method of these instruments allows a note of a given pitch, volume and duration to be triggered at a particular time. Notes are defined using scientific pitch notation (e.g. **F#4**), the notes name (**F#**) combined with the its octave (**4**). Script Lab is able to reference libraries from the Node Package Manager (NPM). Therefore, I was able to test creating notes with pitches defined in the add-in Javascript to confirm Tone.js was suitable for an add-in.

¹<https://github.com/bit101/tones>

²A description of how the note volume changes over its duration

³<https://tonejs.github.io/>

2.2.2 Office Javascript API

In order to create a program for users to produce music from within Excel, the musical output must be informed by the data in the spreadsheet. Tests up to this point had created notes defined within the add-in Javascript. To test the Excel Javascript API, I outputted a note with the Tone.js library, the pitch of which was defined in the spreadsheet. This was extended so the instruction to play a note, not just the pitch, being defined within a cell, detected and executed.

Next, I was able to play a sequence of constant length notes with the notes defined in consecutive cells. The range of cells was accessed using the Excel API and the values were played using the Tone `Sequence` object. Having carried out the above tests, I confirmed Tone.js combined with the Excel API had the functionality required to assist in the implementation of the project.

2.3 Excello Design and Language

2.3.1 Abstracting Time

Dave Griffith's Al-Jazari [7] takes place in a three-dimensional world where robotic agents navigate around a two-dimensional grid. The height and colour of the blocks over which the agents traverse determines the sound that they produce. The characteristics of the blocks are modified manually by users at run-time whilst the agents are moving. Whilst there are more complex conditional instructions, the basic instructions have the agents rotate and move forwards and backwards in the direction that they are facing. There therefore exists a dual formalism as both the instructions given to an agent and the state of each block. This design is intended to make live coding more accessible, both when viewing performances and becoming a live coder.

In Al-Jazari, the agents are programmed by placing symbols corresponding to different movements in thought bubbles that appear above them. This is not suitable for programming within spreadsheets where all data must exist alphanumerically within cells. What's more if an agent was to continue moving forwards many times in a row, it would become tiresome to keep adding the move forward symbol. This is less of an issue in Al-Jazari where the grid within which the agents navigate only measures ten cells wide and long.

The concept of having a cursor navigate around a cartesian plane is the method used by turtle graphics. Just as this concept is used in Al-Jazari to play the cell the agents occupy rather than colour it, it is suitable for spreadsheets. The turtle abstraction is employed by Excello by having notes defined in cells and defining agents, known as turtles, to move through the spreadsheet activating them. In order to play a chord, multiple turtles must be defined to pass through multiple cells corresponding to the note of the chord. This method maintains high notational consistency but sacrifices the abstractions for musical

structures that are available in languages like Sonic Pi - `chord('F#', 'maj7')`. By implementing methods in the add-in to add the notes of chords to the grid, the use of the abstractions is maintained whilst preserving consistency and cleanness in the spreadsheet itself.

The turtle is the crux of the Logo programming language [2]. In Logo, turtles are programmed entirely by text. For example `repeat 4 [forward 50 right 90]` has a turtle move forwards 50 units and turn 90 degrees to the right. This is repeated four times to draw a square. A similar method is employed in Excello but the language is designed to be much more concise.

2.3.2 Initial Prototype Design

In Excello, notes are placed in the cells of the spreadsheet and pathways through the grid are defined using a language for programming turtle movement. The notes in the cells will be played when a turtle moves through that cell. When the program is run, the melodic lines produced by all turtles defined in the grid will be played concurrently. Turtles are defined with a start cell, movement instructions, the speed with which they move through the grid (cells per minute) and the number of times they repeat their path. As in Al-Jazari, distance in space maps to time [8], Excello extends upon this by allowing different turtles to navigate at different speeds. This allows parts with longer notes to be defined more concisely and for phase music⁴ to be easily defined.

As in Logo, turtles begin facing north. The move command `m` moves the turtle forward one cell in the direction that it is facing. Just like in Logo, the turtle always moves in the direction it is facing. The commands `l` and `r` turn the turtle 90 degrees to the left and right respectively. Repeats are implemented in Logo with the command `repeat` followed by the number of repeats and the instructions to be repeated [2]. In order to create more concise instructions, single ommands can be repeated in succession by placing a number immediately after it. For example, the command `m4` will have the turtle move forwards four cells in the direction that it is facing. The direction a turtle is facing can be defined absolutely using the commands `n`, `e`, `s` and `w` to face the turtle north, east, south and west. This could have instead moved the turtle in that direction, but this would have lost the consistency that the turtle always moves in the direction it is facing. In order to change the volume notes are played at, dynamics (`ppp`, `pp`, `p`, `mp`, `mf`, `f`, `ff`, `fff`) can be placed within the turtle instructions. Any notes played after this will be played at that dynamic. In the same way the dynamics in western notation are a property of the staff and not individual notes, dynamics were originally designed to be a property of the turtle. In order to repeat multiple instruction sequences, these are placed in brackets and the number of repeats put immediately after the bracket. For example, `(r m50)4` would define a path going clockwise around a fifty by fifty square. This 8 character example is equivlant to the Logo example above that requires 30 characters. The ability to repeat

⁴Music where identical parts are played at different speeds

Table 2.1: Grammar rules for turtle movement instructions. $z \in \mathbb{Z}, n \in \mathbb{N}, c \in [\mathbf{A-Za-z}]^+$.

Grammar Rule	Left Symbol Meaning
$\mathbf{S} \rightarrow \mathbf{Y}$	Starting symbol
$\mathbf{Y} \rightarrow \mathbf{X} \mathbf{X} \mathbf{Y}$	A series of instructions
$\mathbf{X} \rightarrow (\mathbf{Y})z \mathbf{I}$	A single command or bracketed series of instructions
$\mathbf{I} \rightarrow m z \mathbf{R} \mathbf{R}z \mathbf{A} \mathbf{D} j\mathbf{C} j\mathbf{P}n\mathbf{P}n$	A single command
$\mathbf{R} \rightarrow l r$	Relative rotation
$\mathbf{A} \rightarrow n e s w$	Absolute rotation
$\mathbf{P} \rightarrow + -$	Sign
$\mathbf{C} \rightarrow cz$	Cell reference
$\mathbf{D} \rightarrow ppp pp p mp mf f ff fff$	Dynamic

larger series of instruction is why the relative movements l and r are included in the language despite being less explicit than the compass based directions.

It may not be convenient for each melodic line to be defined by a single path of adjacent cells. Just as conventional score notation often spans across multiple lines, the splitting of parts is a useful form of secondary notation. This requires the turtle to navigate to non-adjacent cells and then proceed playing. For graphic drawing in Logo, the pen can be lifted, allowing the turtle to navigate without colouring the space beneath it. This is suitable for a graphical output where the number of steps the turtle takes has no effect on the output, only the cells it colours. However, the musical output is dependent on when the turtle is in certain cells, so this would not be convenient as it would introduce large rests. Analogous to lifting the pen for graphical turtles, one could set the turtle in a mode where it doesn't play the cells it navigates through and passes through them immediately until it is placed back in a playing mode. However, in this case the actual path that the turtle takes is insignificant only the cell it ends up in. I have therefore added jumps to the language. This can be defined in absolute terms where the destination cell is given (e.g. $jA5$), or relatively (e.g. $j-7+1$), where the number of rows and columns jumped is given instead. An absolute jump may be more explicit to the human reader but defining jumps relatively allows them to be repeated, jumping to different cells in each repeat. For example $r(m7 j-7+1)9 m7$ plays 10 rows of 8 cells from top to bottom playing each row left to right.

The language for turtle movement instructions can be summarised by the following context-free grammar, $(N, \Sigma, S, \mathcal{P})$. Where the non-terminal symbols $N = (\mathbf{S}, \mathbf{Y}, \mathbf{X}, \mathbf{I}, \mathbf{R}, \mathbf{A}, \mathbf{P}, \mathbf{C}, \mathbf{D})$, terminal symbols $\Sigma = (z \in \mathbb{Z}, n \in \mathbb{N}, c \in [\mathbf{A-Za-z}]^+, m, j, l, r, n, e, s, w, +, -, ppp, pp, p, mp, mf, f, ff, fff)$ and starting symbol S . The set of grammar rules are shown in table 2.1:

Notes are defined in the cells using scientific pitch notation - the note name (with accidental⁵ if required) followed by the octave number. Empty cells are interpreted as rests. In order to create notes longer than a single cell, the character s will sustain the note

⁵Sharp or flat symbol used to define a black note on a piano keyboard.

that came before it. This is used to create notes longer than the duration of a single cell. A cell can be sub-divided time-wise into multiple notes by placing multiple notes separated by commas into a cell. The motivation for this design decision was so the length each cell corresponds to is not bound by the length of the smallest note in the piece. For example, a piece defined primarily with crotchets (one unit) but with a single instance of a quaver (half a unit) and dotted crotchet (one and a half units) can define these two notes with `C4,C4` and `s` in two cells. Without this, representing this single quaver would require double the number of cells and introducing many additional `s` cells in the entire piece.

2.4 Software Engineering

2.4.1 Requirements

The success criteria of the project are as follows:

1. Implementation of an API for music playback within a spreadsheet allowing users to:
 - Play individual notes and chords and define their durations.
 - Define multiple parts.
 - Play loops.
 - Define sequences of notes and chords and be able to call these for playback.
 - Define the tempo of playback.
2. Implementation of a converter from MIDI to the spreadsheet representation.
3. Performance of participatory design sessions.
4. Usability testing using participants who have gained familiarity with the system.

In addition to these, the following extension work was completed:

5. Extensions:
 - Implement additional features from issues and requests that arise from participatory design.
 - Explore a compressive conversion from MIDI to the Excel system.

Tools and Technologies Used

Initial tests were written in Javascript in the Script Lab add-in for Excel. Excellio was written in Typescript as this is readily compiled into the Javascript required to run the add-in but provides static type-checking. It also allows the large collection of existing

Javascript libraries to be utilised. Using the Yeoman generator I created a blank Excel add-in project. I used NodeJS to manage dependencies to other Javascript libraries. During development I ran the add-in on a localhost. To allow participants to run Excello on their own machines, I hosted a version of the add-in online using Surge. To run the add-in in Excel, a manifest.xml file is imported which instructs Excel where the add-in is hosted. The converter from MIDI to Excello was implemented in Python using Jupyter Notebooks.

The tone.js library was used to synthesis and schedule sound production via the Web Audio API. The Javascript music theory library tonal was used to produce the notes that make up chords. This prevented the hardcoding of the intervals present in the 109 chords available. The Python library Mido was employed to read python files. All of these libraries have an MIT license.

Starting Point

Having used the Yeoman generator to create an empty Excel add-in, all of the code used to produce Excello and the MIDI converter is produced from scratch using the tools and technologies described above.

I had written simple Javascript for small web pages, but no experience using Node, libraries or building a larger project. I had never used any of the libraries before, therefore, reviewing the documentation was required before and during development. I had gained significant experience with Python and Jupyter Notebooks from a summer internship.

Evaluation Practices

In order to best tune the design of Excello to the needs of potential users, formative evaluation sessions were carried out with participants. As a result, the project followed a spiral development methodology. Due to the number of participants involved and the timeframe of the project, there were only two major development iterations. The first prototype following the design described above, and the second fixing issues and implementing requests brought up by the participants.

The users from the participatory design phase of the project were involved in summative evaluation at the end of the project. By using the same users, I could carry out tests using experienced users of Excello despite the product not yet being released in the public domain.

2.4.2 MIDI files

Musical Instrument Digital Interface (MIDI) details a communications protocol to connect electronic musical instruments with devices for playing, editing and recording music.

A MIDI file consists of event messages describing on/off triggerings for a device or program to control audio [6]. MIDI files were designed to be produced by MIDI controllers such as an electric keyboard. As such, a MIDI file contains a lot of controller specific information that is not necessary for the creation of an Excello file. There exist musical formats such as MusicXML that specify the musical notation and as such may be more suitable for conversion to Excello.

Many musical programs support the importing and exporting of MIDI files. By allowing MIDI files to be converted to the Excello notation, Excello is more integrated into the environment of computer programs for playing, editing and composing music. Furthermore, there exist many datasets available for MIDI [5] which can immediately be played back for comparison.

Chapter 3

Implementation

In this chapter I shall first explain how the turtles, using the movement instructions defined previously, are defined and controlled. The remaining features in the initial prototype will be explained. The format and results of the formative evaluation using this initial prototype shall be summarised. I shall then cover the design decisions and changes that were made to the Excello prototype during the participatory design process. Then, how Excello has been implemented will be described followed by the MIDI to Excello converter. The chapter concludes with an overview of the project repository.

3.1 Initial Prototype

Notes and turtles can be defined in any cell in the spreadsheet. When the Excello add-in is opened, a window will open in the right side of Excel. This contains a play and stop button which can be used to launch all the turtles defined in the spreadsheet and initiate playback. Playback is a realistic piano sound.

A summary of the musical elements that can be put in cells and that will be interpreted by turtles is shown in table 3.1

Table 3.1: Definition of notes in cells.

Interpretation	Format
Note	Note name (A-G), optional accidentals and octave number e.g. F#4
Sustain	s
Multiple notes subdivided in time	Notes, rests or sustains separated by a comma. Rests must be a space or an empty string e.g. E4,,C4,s
Rest	Any cell not interpreted as a note, sustain or multi-note.

3.1.1 Turtles

Notes are played by defining and launching turtles to navigate the spreadsheet. Turtles are defined as follows:

```
!turtle(<Starting Cell>, <Movement>, <Speed>, <Number of Loops>)
```

Activation

The "!" dictates that the turtle will be activated when the play button in the add-in window is pressed. Just like many digital audio workstations allow muting and soloing of tracks, this can be used to quickly modify which turtles will play without losing the definition of the turtle.

Starting Cell

The starting cell of the turtle is given by the cell reference. As with conventional Excel formulae, columns are defined in base 26 using the letters of the alphabet and rows are numbered using the integers. A cell is defined by the concatenation of the column letters and row number. This cell will be played and is the first cell in the path of the turtle.

As each turtle only plays one note at a time, multiple turtles must be defined to play polyphonic music such as chords. It was believed that in these cases the turtles may follow identical paths but in adjacent rows or columns. Multiple turtles following identical paths but starting from adjacent cells can be defined using the existing Excel range notation to define the starting cells. "A2:A5" would define four turtles in the cells A2,A3,A4,A5. This prevents the writing of multiple turtle definitions differing in only the start cell row.

Movement

The turtles start facing north. The design of the language used to define turtle movement has been discussed in the preparation chapter. Using brackets to repeat multiple instructions within the turtle's instructions was not implemented by the start of the participatory design process.

Speed

An optional third argument can be provided to defined the speed at which the turtle moves through the grid. If the argument is not provided, the turtle moves at 160 cells per minute through the grid. The speed argument defines the speed relative to 160 cells per minute. Therefore if the argument "2" was provided, this would move through the grid at 320 cells per minute. This relative system was used so it would be easier to tell the speed relation between two turtles. This would be particularly beneficial for phase music.

Arbitrary maths can be provided for this argument and it will be evaluated. This can be used to define a turtle speed is an irrational multiple of another's.

Number of Loops

An optional fourth argument defines the number of times the path of the turtle is played. If committed the turtle will loop infinitely. This was included so that repeating parts (e.g. the cello part of Bach's Canon in D) need only defining once but can be played repetitively.

3.1.2 Highlighting

In order to assist in the recognition of notes and turtles, when the play button is pressed, cells are highlighted depending on their contents. Cells containing activated or deactivated definitions of turtles are highlighted green. Cells containing definitions of notes, or multiple notes, are highlighted red, with cells only containing a sustain highlighted a lighter red.

3.1.3 Chord input

In order to maintain musical abstractions of chords and arpeggios¹ whilst keeping to paradigm of a turtle being responsible for up to one note at any time, a tool to add chords and arpeggios is available. The note, type (over 100 available), inversion and starting octave of the chord are inputted in four drop-down selectors and the notes making up that chord are entered into the grid using the insert button. If a single cell or range taller than it is wide is highlighted in the spreadsheet, the notes will be inserted in adjacent vertical cells starting at the top-left of the range. Otherwise, the notes will be inserted horizontally. This means whether the turtles are moving horizontally or vertically both chords and arpeggios can be easily defined. As a result, helpful musical abstractions are still available whilst keeping the cleanness of the turtle system.

3.2 Formative Evaluation

In order to guide the development of the project to best suite the users, participants were involved in formative evaluation. 21 participants took place in the participatory design process. Participants were all University of Cambridge students, across a range of subjects, but all musical. Initially, individual meetings were held with each participant. A tutorial of the initial prototype was given followed by the carrying out of a short exercise lasting 15-20 minutes in most cases. After both the tutorial and the exercise, users were

¹Where the notes of a chord are played in rising or descending order

asked to comment on how they found Excello. Particular attention was drawn to actions that they found particularly unintuitive or required notable mental effort. Comparisons were also made to the musical interfaces that participants were already familiar with. The sessions were audio recorded in order to prevent the jotting down of notes causing delays, and later notes were made from these recordings. The ethical and data handling procedures that were carried out shall be discussed in the evaluation chapter.

In order to simulate the most likely ways in which Excello would be used, participants were given the freedom to carry out an exercise of their choice. In many cases this was transcribing an existing piece from memory or from traditional western notation into the Excello notation. Two tasks were provided to choose from if participants had no immediate inspiration. These were to transcribe a piece of music from western notation or to make changes to an existing piece in the Excello notation.

These sessions were carried out at the beginning of Lent term 2019. Participants were asked if they would be willing to continue using Excello personally from their initial session until the summative evaluation sessions, eight weeks later. This gave time for additional feedback to be given as participants used Excello in their own time. It also ensured that the summative evaluation carried out would be done so by users with sufficient experience of the interface. Participants were encouraged to get in contact with any additional issues or suggestions they had during this time.

3.2.1 Issues and Suggestions

The issues and suggestions that arose during the participatory design period have been categorised and summarised below.

Turtle Notation

Having dynamic instructions within the turtle was found to make it harder to extract the path that the turtle followed as not all instructions related to the way in which the turtle moves. As the dynamics weren't next to the notes they corresponded to, it was challenging to establish which volume a given note would be played at or where to place the dynamic instructions within the turtle to correspond to a certain note positioned elsewhere in the spreadsheet. In the initial prototype there was no way to assign a dynamic to the first note without having the starting cell being empty. The addition of this empty cell could be inconvenient for looping parts as this empty cell would be included in the loop. Users who were not familiar with the dynamic markings of western notation found them unintuitive. Furthermore, it was noted that these discrete markings do not make available the continuous volume scale that could be possible with the interface.

When trying to transcribe a piece in an exact tempo, having to divide the speed by 160 in order to enter a relative speed caused unnecessary work. There was forgetfulness as to the whether relative speed referred to how long or how quickly the turtle moved.

Having completed the tutorial, users often has to check which arguments appeared where within the turtle definition or what the arguments were.

When parsing the instruction, as the number of dynamics and movement instructions grew and the instruction became long, it could become quite tough to establish how the turtle would behave. As "s" could be used to indicate sustain within cells, some users confused the "s" within the turtle instructions to mean sustain and not south.

Feedback

Having clicked the play button it was often not clear if this had been registered. Especially if the Excel workbook was saving, and this caused a delay in Excelllo being able to access the spreadsheet. It was also requested if it would be possible to see a summary of where the active turtle were in addition to them being highlighted green. If a turtle had accidentally been left activated, the entire grid had to be searched in order to locate it.

MIDI conversions

Many users, especially those who used production software such as GarageBand and Logic, said that it would be helpful to be able to interact with MIDI files. Being able to use Excelllo to create their chord sequences and bass lines before adding additional effects and recorded lines in their digital audio work stations would require a MIDI export. In addition, if working with an existing MIDI file, it would be convenient to be able to convert that into the Excelllo notation.

Sources of effort when writing

Once notes had been inputted into the grid, often in a single straight line, the number of cells had to be counted so the turtle could be moved the correct number of cells. Whilst Excel allows users to highlight a selection of cells and have an immediate output of how many cells are highlighted without any formula being written, this is still an unproductive use of time. Some users would simply instruct the turtle to move forward significantly more steps than required to prevent this counting, but this is not feasible for looping parts. It was suggested that turtles could figure out how far they should move so that the instruction didn't need editing when more notes were added. This was particularly inconvenient when users were writing out a piece and periodically testing what they had written so far.

If instructions involved repeats such as repeatedly moving to the end of a line and jumping down a few cells and back to the beginning of the line, instructions within the turtles required a lot of repetition.

If writing out a melodic line consisting of small movements, many of the notes would take place in the same octave. As such, it was tiresome to have to repeatedly write out the

octave number when this was barely changing. One user made a comparison to LilyPond² [13] where if the length of a note is not defined, the last defined note length would be used.

Some users said they would find it more intuitive to think of a melodic line in terms of the intervals between notes as opposed to the name of each note. If a piece contained a melodic line that was modulated³, the modulated part had to be written out again and could not be derived quicker from the original version.

Chords

Whilst the variety of available chord types was appreciated, most users used a very small subset of these but still had to scroll through the whole list to find these. Separating the more common chords for easier access was requested. In the initial prototype, the notes were inserted from low to high. This meant that notes inserted in a vertical line had the lowest note at the top with notes increasing in pitch proceeding down the column. In western staff notation, higher pitch notes appear higher up the staff. As a result, it was suggested that inverting the order would be more intuitive. In the initial interface it was also unclear what the different drop downs corresponded to, with some users selecting the "7" from the octave number in order to try and insert a Maj7 chord.

Activation of turtles

When toggling the activation of a turtle, it was very tedious to have to enter the edit mode for each cell containing a turtle definition and add or remove the exclamation mark. This was particularly so if a piece had multiple turtle definitions.

3.3 Second Prototype

Following the formative evaluation sessions carried out with the participants and the feedback that was received, a series of additions and modifications were made to the prototype to solve the problems and opportunities brought up.

3.3.1 Dynamics

In order to assist users in extracting the path that the turtles follow from the instructions and pairing notes with their volume, dynamics are instead inserted in the cells along with the notes. A dynamic instruction is added after the note, separated by a space as in Manhattan [11]. As before, this will persist for all following notes until the volume

²A program for music notation using text notation.

³Where every note has been moved up or down in pitch by the same amount.

is redefined. By moving the dynamics into the cell there is a tradeoff. A single turtle definition with multiple start cells can now play parts of different volume. However, notes in the grid are limited to only being played at the given volume. To play the same notes at a different volume would require a different path to be followed by the turtle where the cells defining the volume are missed and other cells are played. Overall, the new system was believed to be more preferable.

In order to be able to make use of a full continuous dynamic scale, in addition to the existing dynamic symbols available, a number between 0 and 1 can be provided where 0 will be silent and 1 is equivalent to *fff*.

3.3.2 Nested Instructions

The initial language design included the use of nested instruction in order to allow for easy repetition of turtle movements. This was not implemented when the first formative evaluation sessions took place. However, this should help reduce the length of turtle instructions and allow for repeated sections or movements to be more easily incorporated. A series of instructions placed within parentheses with a number immediately following the closing parenthesis will be repeated that number of times. Whilst the fourth argument of the turtle will simply repeat the musical output of the turtle, repetitions within the turtle instruction allow paths to be defined more concisely.

3.3.3 Absolute Tempo

The speed at which the turtle moves is now defined by cells per minute, rather than the relative value used initially. However, values less than 10 were interpreted in the original relative way to maintain backwards compatibility for the participant's existing work. To maintain consistency in a production version, this would be removed so speed must be defined absolutely. This also ensures that the values given for speed and dynamics will be of different orders of magnitude and hence reduce the confusion that can occur between them.

3.3.4 Custom Excel Functions

Two custom Excel functions were implemented to aid in the composition of music within Excel. One to aid with inserting turtles into the grid and a second to transpose notes.

Excello.Turtle

Many users had commented they forgot which arguments go where within the turtle. By adding custom Excel functions the existing formulae writing tools provided within Excel can be utilised. When using a built in formula, a prompt appears informing users which

arguments go where and whether they are optional. The output of this function is text used to define a turtle if written manually. This also allows other cells to be referenced for the arguments of the turtle function. For example a cell could have speed defined that all other turtles reference. As a result, the speed of all turtles could be modified by changing this single value. This also allows relative tempos to be easily implemented again as the speed argument of the turtle could be defined as a relative speed multiplied by the global speed (given as a cell reference).

Excello.Modulate

The implementation of a function to modulate notes both allows for the easy modulation of existing sections of a piece and also the definition of a melodic line by the intervals between the notes. The function takes the contents of a cell where a note is defined and an interval and outputs the note transposed by that interval. A section can be modulated by calling this function on the first note with a provided interval and using the existing drag-fill functionality of Excel to modulate all notes. By using the previous note that has just been transposed and one of a series of intervals as the arguments, a melodic line can quickly be produced from a starting note and a series of intervals.

??

3.3.5 Sustain

In order to prevent confusion between the instruction for a turtle to face south and for a note to be sustained. The symbol “-” has been chosen to sustain a note. This was chosen because it is light and also has some similarity to a tie⁴. The use of an “s” is still interpreted as a sustain to maintain backwards compatibility for the existing work of the participants.

3.3.6 Active Turtles

In order to provide feedback that turtle definitions have been recognised, in addition to green highlighting, a list of the active turtles is given below the play button. This also assists in finding any spurious turtles that were not intended to be activated.

3.3.7 Automatic Movement

To prevent the number of cells in a line needing counting, a turtle can be instructed to move as far as there are notes defined in the direction it is currently facing. This means that if more notes are added on this line, the turtle instructions do not need editing before pressing play. There may be cases where a part is meant to finish with a number of rests.

⁴A line to increase the length of a note by joining to another.

As a rest is notated with a blank cell, a method of increasing the length of the path to include these rests is required. A rest can be given explicitly with a "." allowing the distance distance travelled by the turtle when moving automatically to be increased. This would be required if multiple turtles were defining a repeating section where one does not have the final cell of the section being a note, sustain or multi-note cell. Without an explicit rest the turtle would stop and repeat too soon and the parts would be out of phase.

3.3.8 Inferred Octave

To prevent the octave number needing writing every time a note is defined, the octave number can be inferred by the program if it is omitted. There were two methods under consideration. Firstly, given that most intervals within music are small, the nearest note could be inferred. This means that a scale would only need the octave defining in the first note. Whilst this method would likely require the least explicit statement of octave number it would be non-trivial to figure out the octave a given note is played in. The last defined octave in the path would need finding and then the subsequent notes would have to be walked through keeping track of which octave is being played. The second consideration was to always use the last defined octave. Whilst this may require reasonable octave definition around the boundary between octaves, it is easier to find what octave a note is played at as it is simply the last defined octave in the path.

3.3.9 Chords

In order to assist in the entering of common chords, common chord types are repeated in a separate group at the top of the type selection drop-down. The layout of the chord drop-downs was improved with labels added to make it clearer what the different values would refer to. If the notes were entered vertically, the order was reversed to have a greater correspondence with traditional staff notation.

3.3.10 Activation of turtles

A "Toggle Activation" button was added to the add-in window. When a cell or range is highlighted in the spreadsheet, the activation of any turtle definitions in this range will be toggled when the button is pressed. This significantly increases the ease with which turtles can be deselected as only two clicks are required as opposed to having to enter the cell edit mode and add or remove an exclamation mark.

3.4 How it works

3.5 MIDI Converter

3.6 Repository Overview

Bibliography

- [1] Samuel Aaron, Alan F. Blackwell, and Pamela Burnard. The development of sonic pi and its use in educational partnerships: Co-creating pedagogies for learning computer programming. *Journal of Music, Technology and Education*, 9:75–94, 05 2016.
- [2] Ron Goldman, Scott Schaefer, and Tao Ju. Turtle geometry in computer graphics and computer-aided design. *Computer-Aided Design*, 36:1471–1482, 2004.
- [3] Thomas Green and Alan Blackwell. Cognitive dimensions of information artefacts: a tutorial. Technical Report Version 1.2, BCS HCI Conference, 1998.
- [4] Sven Gregori. Never mind the sheet music, heres spreadsheet music, 2019.
- [5] Allen Huang and Raymond Wu. Deep learning for music. *CoRR*, abs/1606.04930, 2016.
- [6] D.M. Huber. *The MIDI Manual: A Practical Guide to MIDI in the Project Studio*. Taylor & Francis, 2012.
- [7] Alex Mclean, Dave Griffiths, Foam Vzw, Dave@fo Am, Nick Collins, and Geraint Wiggins. Visualisation of live code. 01 2010.
- [8] Alex Mclean and Geraint Wiggins. Texture: Visual notation for live coding of pattern. 01 2011.
- [9] Mozilla. Web audio api. https://developer.mozilla.org/en-US/docs/Web/API/Web_Audio_API, 03 2019. Accessed: 2019-04-02).
- [10] Michael Muller and Sarah Kuhn. Participatory design. *Communications of the ACM*, 36:24–28, 06 1993.
- [11] Chris Nash. Manhattan: End-user programming for music. In *NIME*, 2014.
- [12] Simon Peyton Jones, Margaret Burnett, and Alan Blackwell. A user-centred approach to functions in excel. June 2003.
- [13] Erik Sandberg, Examensarbete Nv, Reviewer Arne Andersson, and Examiner Anders Jansson. Separating input language and formatter in gnu lilypond, 2006.
- [14] Advait Sarkar. Towards spreadsheet tools for end-user music programming. In *PPIG*, 2016.

- [15] C. Scaffidi, M. Shaw, and B. Myers. Estimating the numbers of end users and end user programmers. In *2005 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC'05)*, pages 207–214, Sep. 2005.