

# Learned Indexes

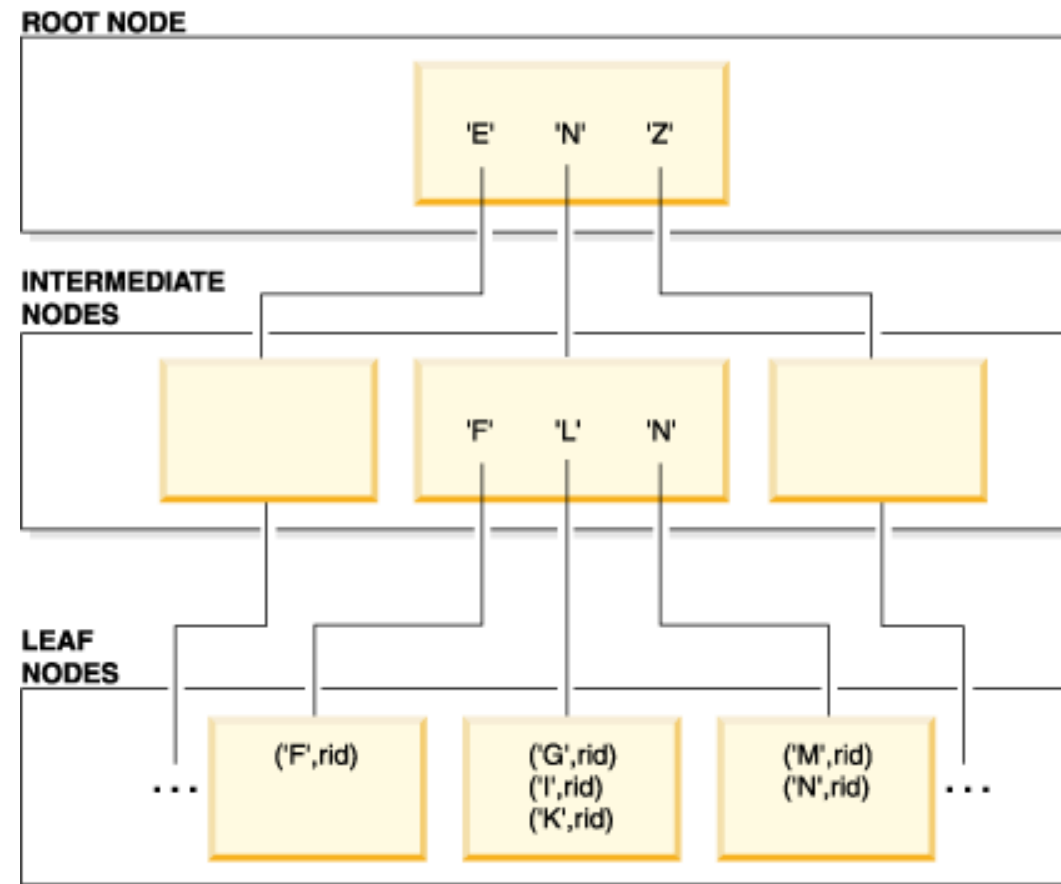
Part II Project - Franz Nowak

# What's a database index again?

- data structure that improves speed of data retrieval operations
- cost of additional writes
- cost of storage space to maintain index
- e.g. b-tree on sorted data: lookup  $O(\log(n))$  -> optimal for range lookups

# What's a database index again?

b-tree index




source:<https://www.ibm.com/support/knowledgecenter>


# Exploit inherent structure of data

- What if data is all integers from 0 to 1,000,000?
- `data_array[lookup_key]`
- No need for b-tree, lookup is  $O(1)$  time and space!



# How can we improve on this?



# Core idea: B-Trees are models



# Key idea: B-Trees are models



# To Do: Regression Tree

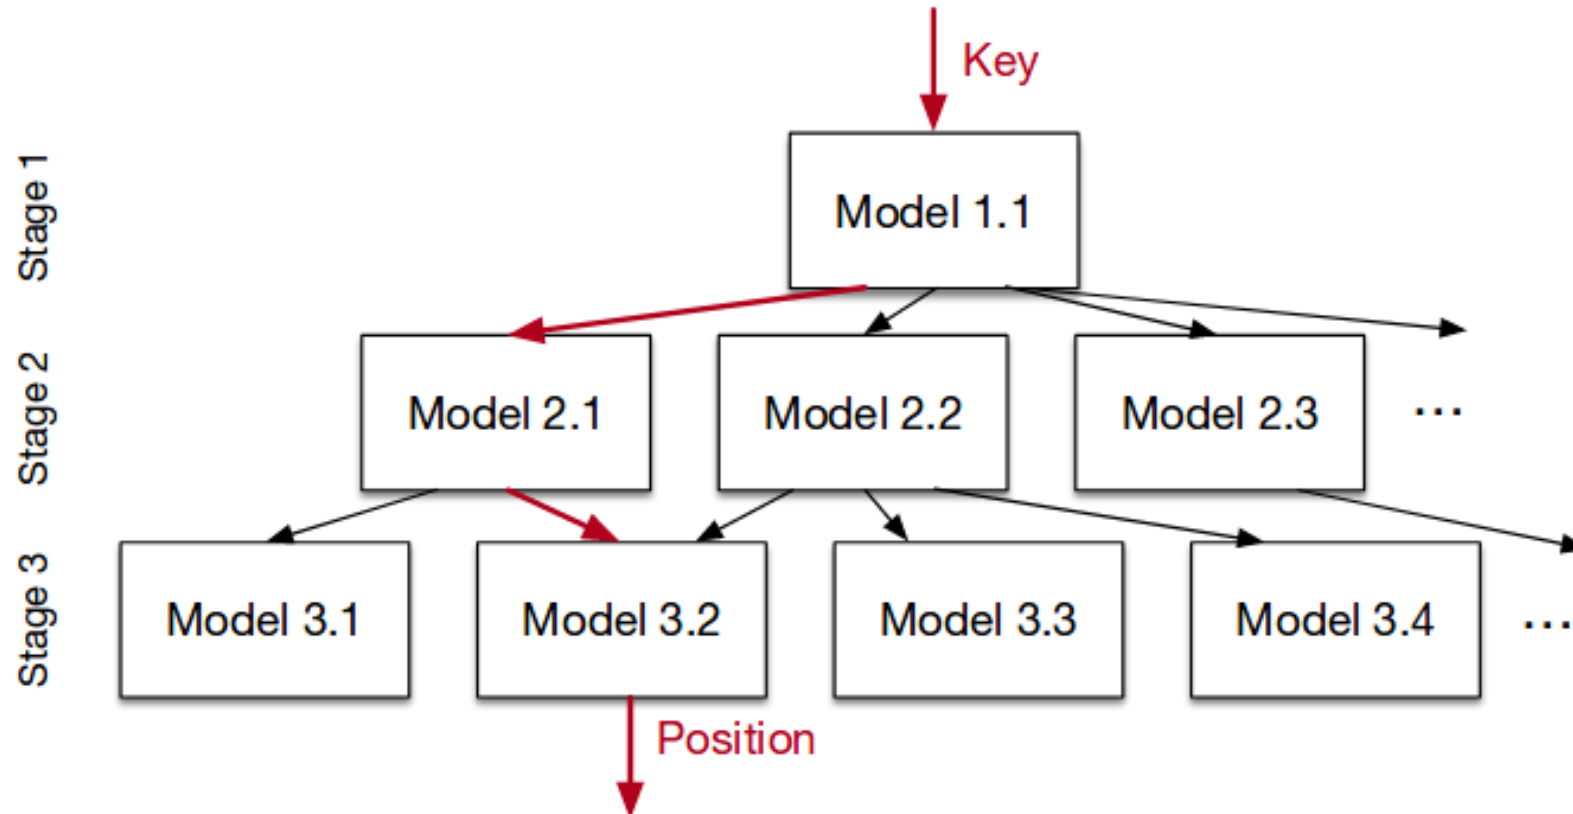


Figure 3: Staged models



# Position in sorted array: CDF


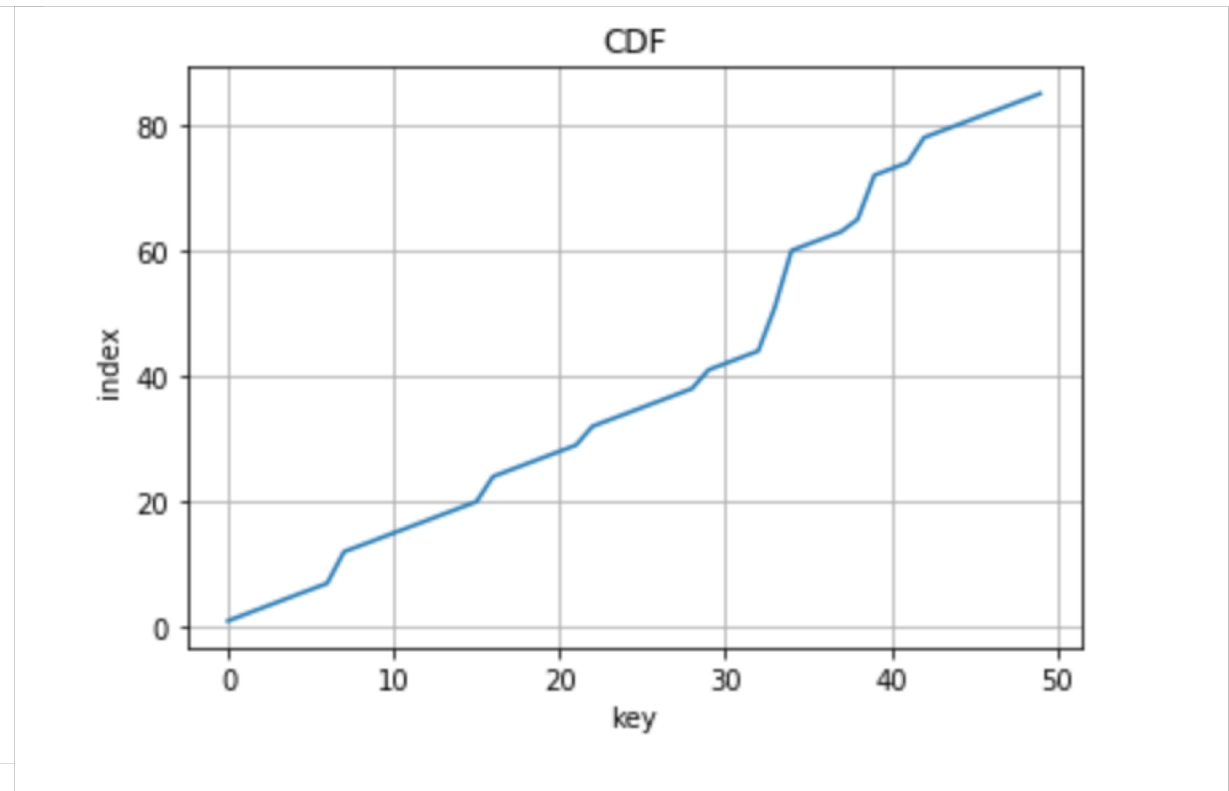
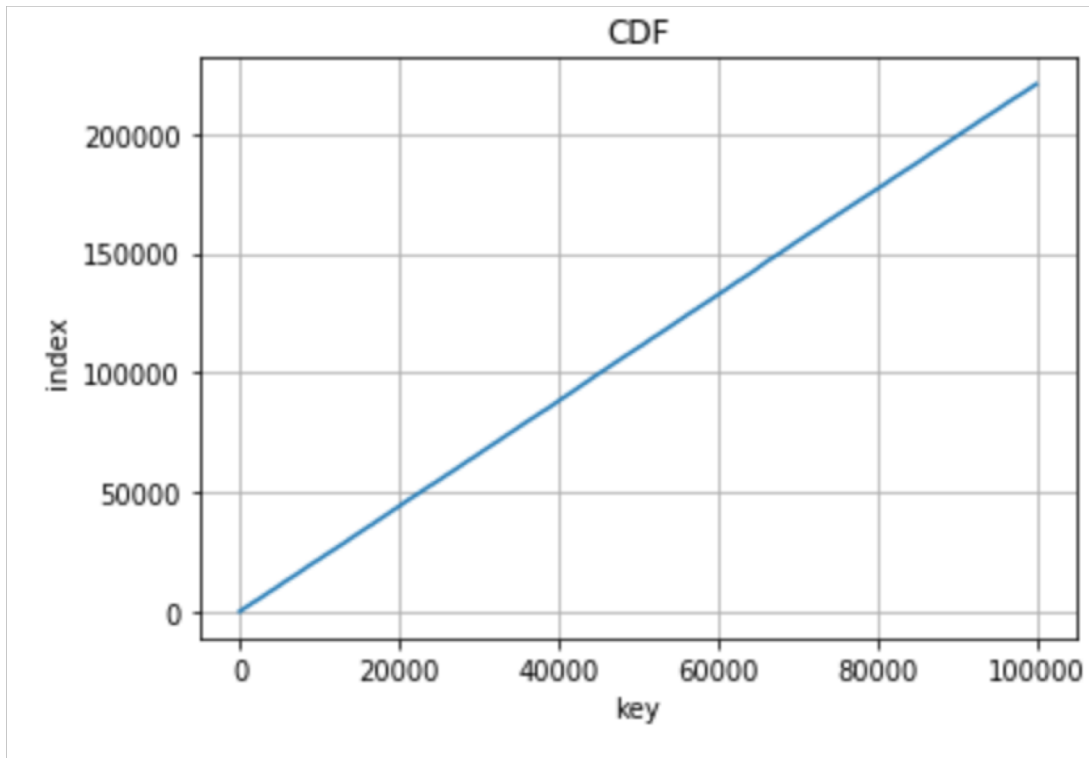


Figure 2: Indexes as CDFs

# Best case vs worst case




# Dataset

- 100,000 keys in sorted array
- 10 levels of entropy
- averaged over 100 runs

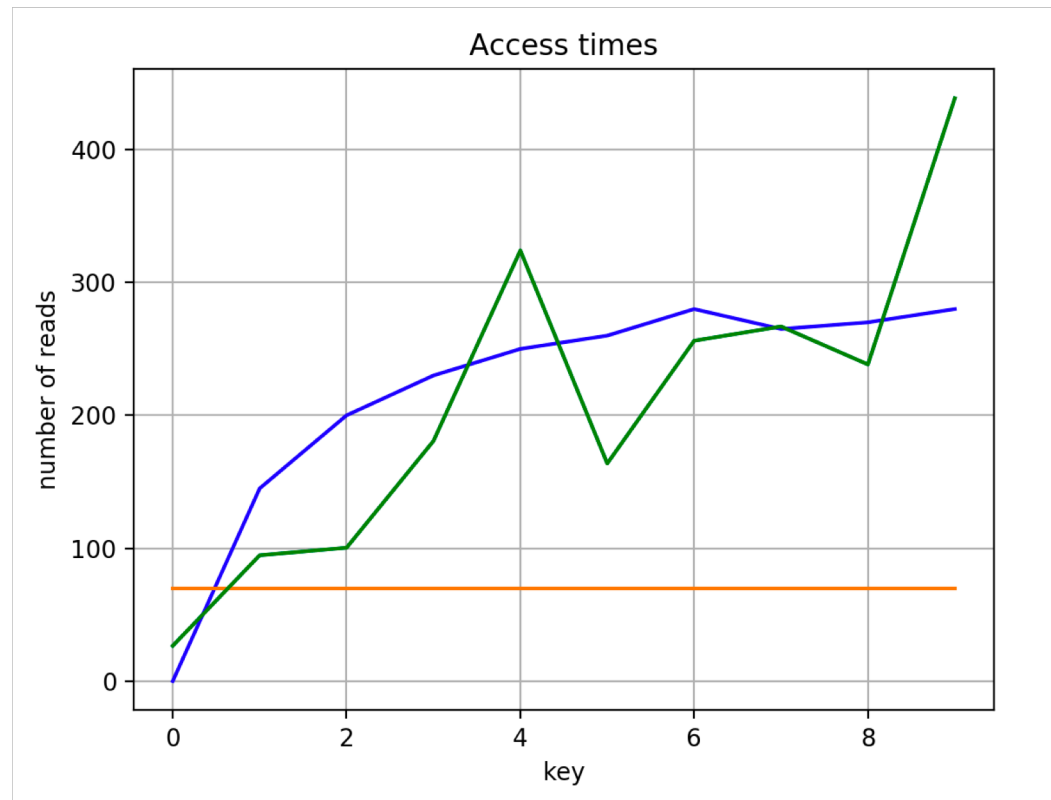
# Implementation-independent evaluation

- What we really want to know: wall-clock time
- different depending on OS, hardware (SIMD), implementation of b-tree, programming language, ...
- Idea: use reads as proxy
- $\text{read\_time} = f(n) + \text{reads}$
- $f(n)$  is the base complexity of the data structure (e.g. invocation overhead of NN of size proportional to dataset size)


# Preliminary results



# Preliminary results



# Hypothesis



# Main references

- The paper: <https://arxiv.org/abs/1712.01208>
- Stanford EE380: Computer Systems Colloquium Seminar
- <https://www.youtube.com/watch?v=NaqJO7rrXy0>